

JOINT EVENT LOCALIZATION AND CAPTIONING USING PRETRAINED SEQUENCE-TO-SEQUENCE MODELS FOR DENSE VIDEO CAPTIONING

Huynh Ngoc Bao Long

Nguyen Huu Sang

University of Information Technology - National University of Vietnam

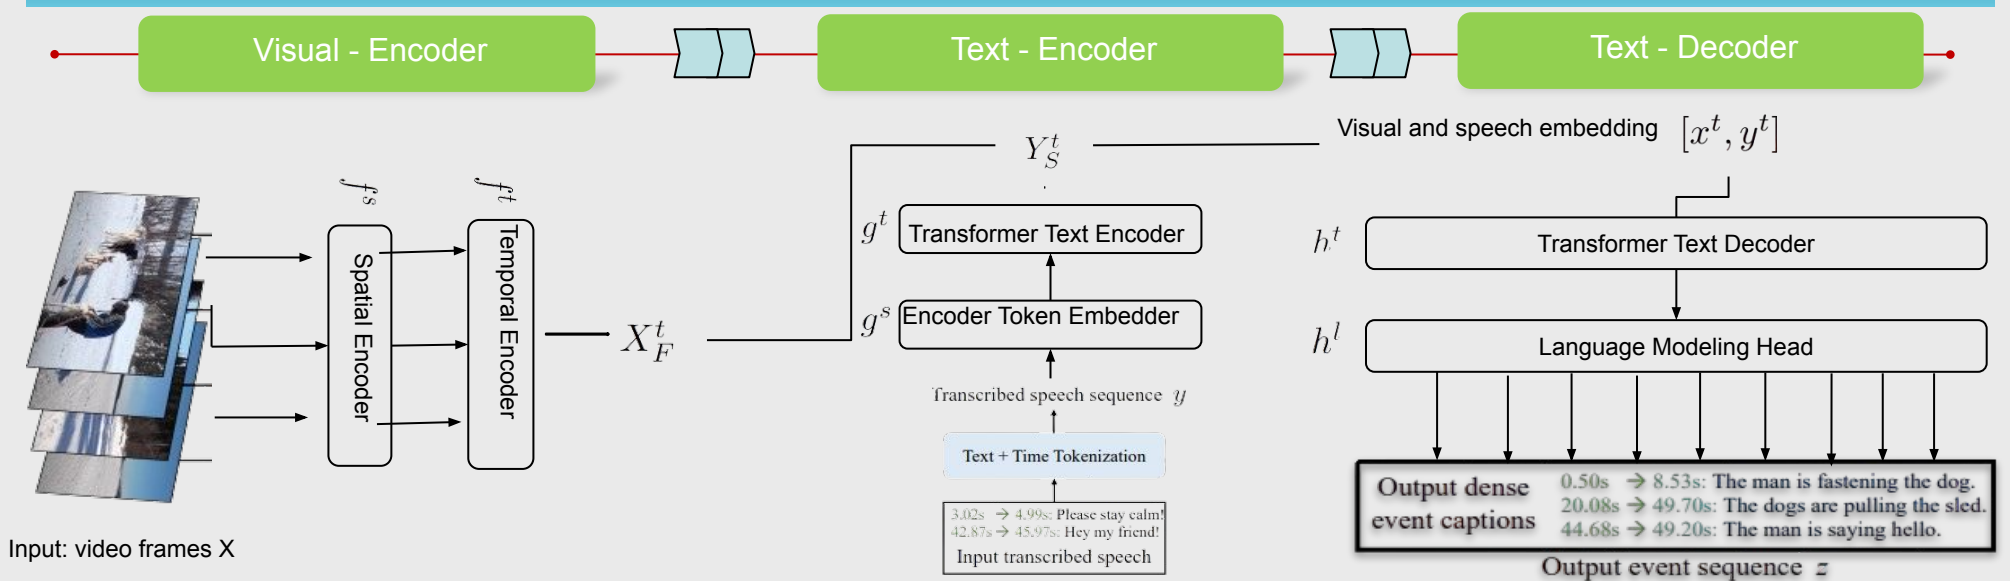
Motivations

The rapid growth of video content on digital platforms has made automatic video understanding a critical challenge. **Dense video captioning**, which identifies key events in untrimmed videos and generates concise descriptions, is crucial for multimedia analysis, accessibility, security, and data management. However, existing **two-stage methods** are inefficient and rely heavily on manual annotations. This research proposes a unified system integrating **event localization** and **captioning** using pretrained **sequence-to-sequence** deep learning models. By leveraging unannotated data and simplifying the processing pipeline, the approach improves scalability and performance, offering practical solutions for real-world applications and advancing video and natural language processing.

Targets

- **Develop an integrated sequence-to-sequence model:** Build a state-of-the-art model capable of simultaneously localizing and generating detailed, temporally coherent descriptions of events in videos.
- **Leverage video data with transcripts:** Utilize natural, unannotated video data to reduce costs while effectively capturing relationships between visuals, audio, and language.
- **Evaluate fine-tuning performance:** Measure model effectiveness through log-likelihood, assessing its predictive accuracy and the quality of generated descriptions compared to real-world data.

Overview



Description

1. Pre-training Dataset

- We utilized a large-scale dataset comprising narrated videos, offering an extensive collection of video-text pairs.
- This dataset encompasses diverse scenarios and natural narrations, providing valuable data for training models to produce detailed and coherent video descriptions.

- **Denoising objective:** Reconstruct masked tokens by leveraging both noisy speech and visual inputs, promoting the model's ability to reason across modalities and handle incomplete or noisy data effectively.

$$\mathcal{L}_{\theta}(x, y, z) = -\frac{1}{\sum_{k=1}^{L-1} w_k} \sum_{k=1}^{L-1} w_k \log p_{\theta}(z_{k+1} | x, y, z_{1:k})$$

Equation 1 . Likelihood loss function

4. Research Plan

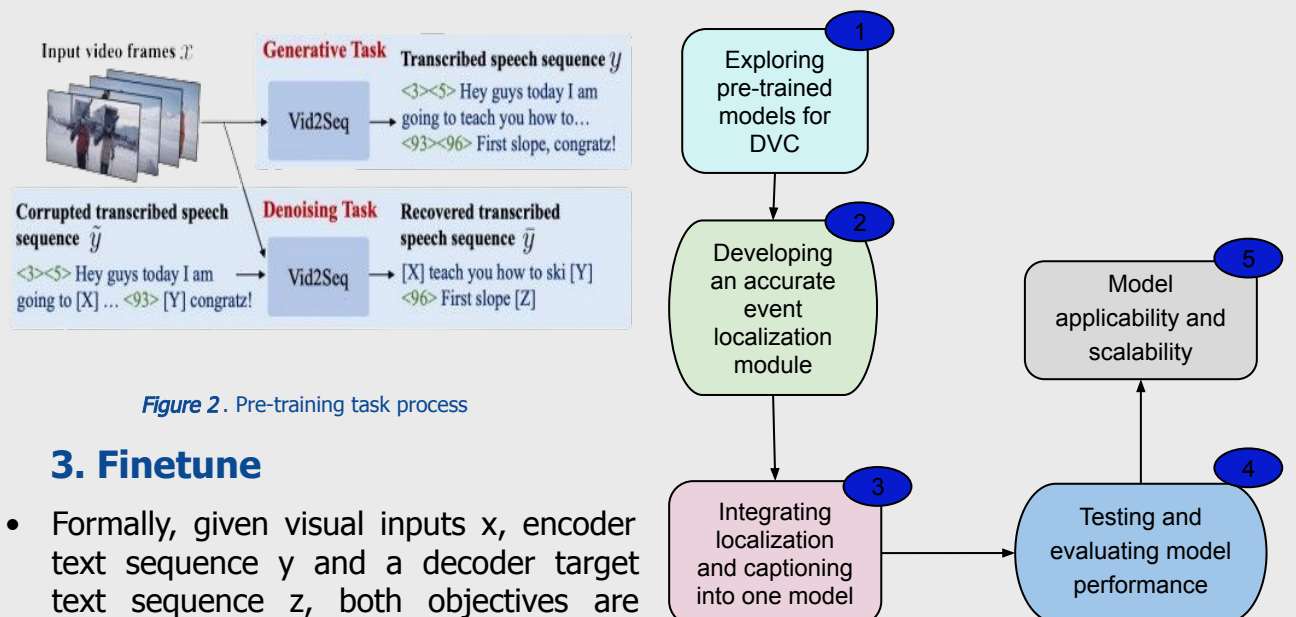


Figure 3 . Research plan diagram

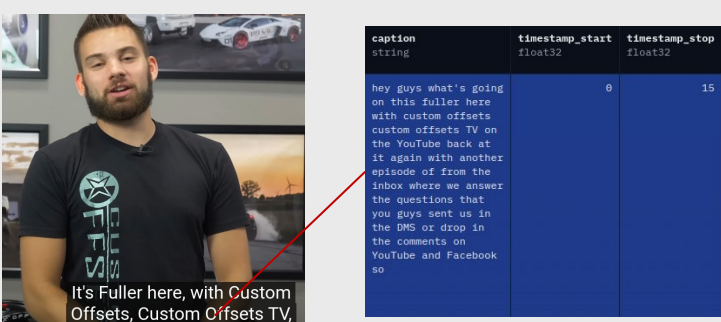


Figure 1 . From youtube subtitles to timestamp data

2. Pre Training Task

- **Generative objective:** Generate coherent speech outputs conditioned on visual inputs, enabling the model to learn a strong connection between visual and linguistic modalities.

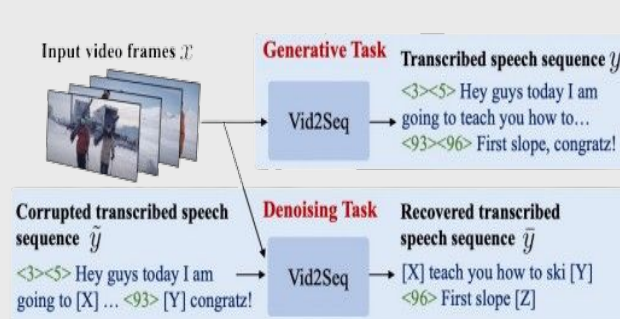


Figure 2 . Pre-training task process

3. Finetune

- Formally, given visual inputs x , encoder text sequence y and a decoder target text sequence z , both objectives are based on minimizing the following loss: